

## Database Systems

### Algorithms for Relational Algebra Operators and Query Evaluation

**Dr P Sreenivasa Kumar**

Professor

CS&E Department

IIT Madras

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

1

---

---

---

---

---

---

---

---

## Relational Query Evaluation

- Relational Algebra Operators
  - Select, Project, Join
  - Union, Intersect, Difference
- Grouping and aggregation
  - Sorting
- How to implement these?
- How do indexes help?
- Any other information is helpful?

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

2

---

---

---

---

---

---

---

---

## Selection With Equality Conditions

- Single selection condition  $X = c_i$ 
  - Index on  $X$  ? Yes: use the index; No: file scan
- Several conjunctive conditions
  - $X_1 = c_1$  and  $X_2 = c_2$  and ... and  $X_k = c_k$ 
    - Index on any  $X_i$  ?
      - Yes: Get the records and check other conditions
      - No: File scan
- Several disjunctive conditions
  - Index on any *single*  $X_i$  - not helpful
    - Difficult compared to conjunctive case

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

3

---

---

---

---

---

---

---

---

## Predicate Selectivity

- Selectivity  $s$  of a condition  $C$  --  $0 \leq s \leq 1$ 
  - (No. of records satisfying  $C$ ) / (Total no. of records)
    - $C_1$ : student.dept = "CSE" --  $450 / 8000 = 0.056$
    - $C_2$ : student.sex = "female" --  $1200 / 8000 = 0.15$
    - $C_3$ : student.rollNo = "CS10B032" --  $1/8000 = 0.000125$
  - highly selective predicate - very **low** selectivity value
- Conjunction of conditions
  - Choose the one that is *most* selective
    - Get the records and check other conditions
  - Selectivity values (estimates): collect offline

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

4

## Selectivity Estimation

- Maintained in the DB *catalog*
  - Used by the query optimizer
- Equality conditions involving a key attribute
  - Selectivity =  $1 / (\text{Total no. of records})$
- Equality conditions involving a non-key attribute
  - Selectivity =  $1 / (\text{Distinct values of the attribute})$
- Sometimes histograms are also maintained
  - Distinct value or value range -- # of records

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

5

## Project Operation

- For every record in the operand
  - Access it, take the required attributes values
  - Construct the result record
- Duplicate Elimination
  - Costly
  - Sort or hash based methods are used
- File scan becomes essential
- Apply project after selection, if possible
  - To reduce the input to project

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

6

## External Sorting

- Sorting a file
  - An often required operation
    - Duplicate elimination, Grouping of records, Join etc
- Merge-sort Principle is used
  - $O(n \log n)$  worst-case complexity for  $n$  items
  - Two phases
    - **Sort phase** – repeat: read part of data, sort and write
      - Create many sorted files – called *runs*
    - **Merge phase** – repeat: merge *some* sorted files and write
      - Till only one sorted file is left

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

7

---

---

---

---

---

---

---

---

## Algorithm – Sort Phase

- File:  $n$  blocks and Buffer memory:  $m$  blocks
- Sort Phase
  - Repeat the following  $\lceil n/m \rceil$  times
    - {read the next  $m$  blocks; sort in-memory;  
write to disk as a single file, called a *run*}
- Number of *runs*  $r = \lceil n/m \rceil$
- Complexity:  $n$  block reads and  $n$  block writes
  - $2n$  block accesses

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

8

---

---

---

---

---

---

---

---

## Algorithm – Merge Phase

- File:  $n$  blocks, Memory Buffers:  $m (\geq 3)$  blocks, Runs:  $r$ 
  - Degree of merging  $d$ :  $2 \leq d \leq (m-1)$
- Merge Phase: repeat the following  $\lceil \log_d r \rceil$  times
  - Reduce  $j$  runs to  $\lceil j/d \rceil$  runs (Initially,  $j = r$ )
    - By repeatedly merging  $d$  runs at a time to get *one* run
      - Use  $d$  buffers, one for each of the next  $d$  runs; use one for the result
      - Get one block at a time from each run
      - Merge and write the result to disk – one block at a time
- Complexity:  $2n \lceil \log_d r \rceil$ 
  - Each sub-phase : Entire file gets read and written
- Overall:  $(2n + 2n \lceil \log_d r \rceil)$  block accesses

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

9

---

---

---

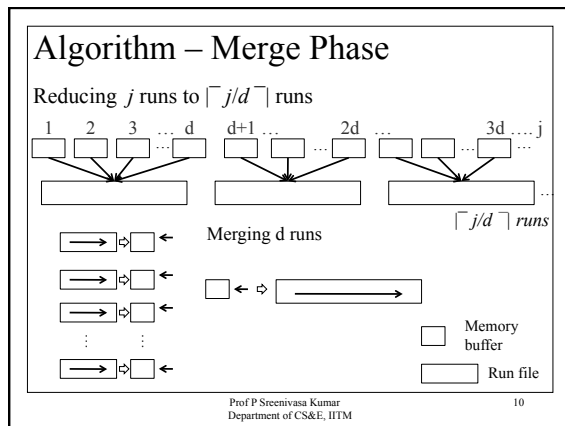
---

---

---

---

---




---

---

---

---

---

---

---

---

### Join Processing

- Join – A very important operation
- 2-way join
  - Two files of records, join condition – given
- Multi-way join
- Choice of algorithm depends on ...
  - Sizes of files
  - Primary organization of the files
  - Availability of indices
  - Selectivity of the join condition etc

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

11

---

---

---

---

---

---

---

---

### Nested Loop Join (or block nested loop join)

- Brute force join
- Two data files
 

for each record  $x$  in  $R$  do  
 for each record  $y$  in  $S$  do  
 check if  $x, y$  join ...
- $R : b_1$  blocks,  $S : b_2$  blocks, Buffer :  $m$  blocks
- Buffer Usage: One block for the result of join
  - One for inner file (say,  $S$ );  $(m - 2)$  for outer file ( $R$ )
- For each set of  $(m - 2)$  blocks of  $R$  read-in, do
  - For each block of  $S$  do
    - Read it in, compute join, write to result block
    - Write the result block to disk whenever it fills up

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

12

---

---

---

---

---

---

---

---

### Nested Loop Join - Performance

- Two data files
  - $R : b_1$  blocks,  $S : b_2$  blocks, Buffer :  $m$  blocks
- Outer file :  $b_1$  blocks accesses
- # times inner file blocks accessed:  $\lceil b_1/(m-2) \rceil$
- Overall:  $b_1 + \lceil b_1/(m-2) \rceil b_2$
- Or, symmetrically:  $b_2 + \lceil b_2/(m-2) \rceil b_1$ 
  - when we have S in the outer loop and R inside
- Which file in the outer loop?
 

Time for writing the result needs to be added

  - The one with fewer blocks!

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

13

### Nested Loop Join - Example

- Two data files
  - $R : b_1 = 5600$  blocks,  $S : b_2 = 120$  blocks, Buffer : 52 blocks
- If R is used in the outer loop
  - $b_1 + \lceil b_1/(m-2) \rceil b_2$
  - $5600 + \lceil 5600 / 50 \rceil * 120 = 19040$  disk ops
- If S is used in the outer loop
  - $120 + \lceil 120 / 50 \rceil * 5600 = 16920$  disk ops
- Assuming 10 msec per disk op
 

Time for writing the result needs to be added

  - It is 190 secs versus 169 secs

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

14

### Single Loop Join (or index loop join)

- Two data files
 

Time for writing the result needs to be added

  - $R : b_1$  blocks,  $S : b_2$  blocks
    - Need to compute **equi-join** with  $R.A = S.B$
  - We have index on one of them, say S on B
- For each record  $x$  of R read in, do
  - Use the index on B for S
  - Get all the matching records (having  $B = x.A$ )
- Time taken:  $b_1 + |\text{distinct}(R.A)| * h_B(S)$ 
  - $h_B(S)$  – # of block accesses of the index on B for S

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

15

### Join Selection Factor

- Fraction of records in a file that join with records of the other for the given condition
- Consider: professor  $\bowtie_{\text{empId} = \text{hod}}$  department
  - Only 5% of professor rows join with department rows
  - 100% of department rows join with professor rows
- Impacts performance of single loop join
  - If indexes are available on both files
  - Loop over records of the file with high join selection factor

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

16

---

---

---

---

---

---

---

---

### Join Selection Factor - Example

- Impacts single loop join performance
  - If indexes are available on both files
- Consider: professor  $\bowtie_{\text{empId} = \text{hod}}$  department
  - Loop over *professor* records and probe *department* using index on *hod* (option 1) OR
  - Loop over *department* records and probe *professor* using index on *empId* (option 2)
  - Option 1: 95% probes don't give a match
  - Option 2: All probes give a match
- Option 2 is the right choice

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

17

---

---

---

---

---

---

---

---

### Hash Join

- Consider a 2-way equi-join  $R \bowtie_{R.A=S.B} S$ 
  - Assume that *S* fits into memory
- Use a hash function *h*
  - Hash the records of *S* into *M* buckets using *B*-values
    - Called the **partitioning** of *S*
- To compute join result
  - Hash records of *R*, one by one, using *A* values
    - Use the *same* *M* buckets and the *same* hash function *h*
    - Matching pair of records will hash to same bucket

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

18

---

---

---

---

---

---

---

---

## Partition Hash Join

- Consider a 2-way equi-join  $R \bowtie_{R.A=S.B} S$ 
  - **Neither R nor S fits into the memory**
- Partition Phase: use a hash function  $h$ 
  - Hash the records of R into  $m$  buckets using A-values
    - We get  $R_1, R_2, \dots, R_m$  - write them to files
  - Hash the records of S into  $m$  buckets using B-values
    - We get  $S_1, S_2, \dots, S_m$  - write them to files
  - Goals: ensure that distribution is uniform and
    - At least one of  $R_i$  or  $S_i$  fit into the memory
- To compute join result: join  $R_i$  with  $S_i$  only!

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

19

---

---

---

---

---

---

---

---

## Partition Hash Join – Probe Phase

- Probe Phase: Join  $R_i$  with  $S_i$  for all  $i$
- If one of  $R_i$  or  $S_i$  fit into the memory
  - Use the idea of hash join again!
    - Hash the smaller of the two into main memory using a *different* hash function, say  $h_2$
    - Read the other file, probe and produce result records
  - Overall cost:  $(3(|R|+|S|) + |\text{result}|)$  block accesses
- Else use nested loops join
  - Overall cost:  $2(|R|+|S|) + \text{cost of nested loop joins}$

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

20

---

---

---

---

---

---

---

---

## Sort-merge join

- Consider a 2-way equi-join  $R \bowtie_{R.A=S.B} S$
- If R is sorted on A, S is sorted on B
  - Merge R and S to get join results
  - Called merge join - - very efficient - - linear
- If one of them is sorted on join attribute
  - Sorting the other and merging may be cost-effective
- Of course, we can
  - Sort R on A, sort S on B and use merge
  - Cost might be high

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

21

---

---

---

---

---

---

---

---

## Set Operations

- Hash based join method
  - Can be adapted to compute Union, Intersect and Difference
- Sort-Merge method
  - Can be adapted to compute Union, Intersect and Difference
- Please study the details!

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

22

---

---

---

---

---

---

---

---

## Query Optimization

- An SQL query - converted to a RA expression tree
- Initial RA expression is re-written
  - Using heuristic and algebraic transformation rules that preserve the meaning of the expression
    - Called algebraic optimization
  - Final RA expression tree is generated
- Cost-based query optimization
  - Cost estimates of *methods* for RA ops are computed
  - Execution plan with least estimated cost is chosen

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

23

---

---

---

---

---

---

---

---

## Heuristic Optimization

- An SQL query - converted to a RA expression tree
- This RA expression tree is to be re-written
- Main heuristic rule
  - Apply *select* and *project* before other operations
    - Reduces the size of intermediate results
    - Reduces the number of fields in the intermediate results
- Make use of relational algebraic laws
  - *Select, project, join, union, intersect* - commutative
  - *Join, union, intersect* - associative
  - There are many more....(Read about them)

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

24

---

---

---

---

---

---

---

---



## Cost-based Optimization

- After initial RA expression tree is re-written using heuristics and algebraic laws....
- Each RA operator
  - Can be evaluated using *many* methods
  - For a method, its *cost function* gives *estimated cost*
    - By taking file sizes, access path costs etc into account
  - Choice made at a node may effect choices at others
- Evaluate different plans based on estimated costs
  - Choose the plan with least estimated cost

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

25

---

---

---

---

---

---

---

---

## Query Optimization – Example

- Obtain the name and phone details of professors who taught the courses taken by student with roll number “CS08B027” in the even semester of 2010
- *select* p.empId, p.name, p.phone  
*from* professor p, teaching t, enrollment e  
*where* e.rollNo = “CS08B027”  
and e.courseId = t.courseId  
and e.sem = “even” and e.year = 2010  
and t.sem = “even” and t.year = 2010  
and p.empId = t.empId

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

26

---

---

---

---

---

---

---

---

## Query Optimization – Example

- Obtain the name and phone details of professors who taught the courses taken by student with roll number “CS08B027” in the even semester of 2010
- Initial RA Expr:  $\Pi_{p.empId, p.name, p.phone} ( \sigma_{\theta} (p \times t \times e) )$   
*where*  
p: professor, t: teaching, e: enrollment  
 $\theta = (e.rollNo = \text{“CS08B027” and } e.courseId = t.courseId$   
and e.sem = “even” and e.year = 2010  
and t.sem = “even” and t.year = 2010  
and p.empId = t.empId)

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

27

---

---

---

---

---

---

---

---

### Query Optimization – Example

$$\begin{aligned} & \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta} ( p \times t \times e ) ) \\ & \equiv \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta_3} ( p \times \sigma_{\theta_2} ( t ) \times \sigma_{\theta_1} ( e ) ) ) \end{aligned}$$

p: professor, t: teaching, e: enrollment

$\theta_1 = (e.\text{rollNo} = \text{"CS08B027"} \text{ and } e.\text{sem} = \text{"even"} \text{ and } e.\text{year} = 2010)$

$\theta_2 = (t.\text{sem} = \text{"even"} \text{ and } t.\text{year} = 2010)$

$\theta_3 = (p.\text{empId} = t.\text{empId} \text{ and } e.\text{courseId} = t.\text{courseId})$

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

28

### Query Optimization – Example

$$\begin{aligned} & \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta} ( p \times t \times e ) ) \\ & \equiv \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta_3} ( p \times \sigma_{\theta_2} ( t ) \times \sigma_{\theta_1} ( e ) ) ) \\ & \equiv \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta_3} ( p \times \sigma_{\theta_4} ( \sigma_{\theta_2} ( t ) \times \sigma_{\theta_1} ( e ) ) ) ) \end{aligned}$$

p: professor, t: teaching, e: enrollment

$\theta_1 = (e.\text{rollNo} = \text{"CS08B027"} \text{ and } e.\text{sem} = \text{"even"} \text{ and } e.\text{year} = 2010)$

$\theta_2 = (t.\text{sem} = \text{"even"} \text{ and } t.\text{year} = 2010)$

$\theta_3 = (p.\text{empId} = t.\text{empId})$

$\theta_4 = (e.\text{courseId} = t.\text{courseId})$

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

29

### Query Optimization – Example

$$\begin{aligned} & \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta} ( p \times t \times e ) ) \\ & \equiv \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta_3} ( p \times \sigma_{\theta_2} ( t ) \times \sigma_{\theta_1} ( e ) ) ) \\ & \equiv \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta_3} ( p \times \sigma_{\theta_4} ( \sigma_{\theta_2} ( t ) \times \sigma_{\theta_1} ( e ) ) ) ) \\ & \equiv \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( p \stackrel{\sim}{\bowtie}_{\theta_3} ( \sigma_{\theta_2} ( t ) \stackrel{\sim}{\bowtie}_{\theta_4} \sigma_{\theta_1} ( e ) ) ) \end{aligned}$$

$\theta_1 = (e.\text{rollNo} = \text{"CS08B027"} \text{ and } e.\text{sem} = \text{"even"} \text{ and } e.\text{year} = 2010)$

$\theta_2 = (t.\text{sem} = \text{"even"} \text{ and } t.\text{year} = 2010)$

$\theta_3 = (p.\text{empId} = t.\text{empId})$

$\theta_4 = (e.\text{courseId} = t.\text{courseId})$

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

30

## Query Optimization – Example

$$\begin{aligned}
 & \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta} ( p \times t \times e ) ) \\
 \equiv & \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta_3} ( p \times \sigma_{\theta_2}(t) \times \sigma_{\theta_1}(e) ) ) \\
 \equiv & \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta_3} ( p \times \sigma_{\theta_4} ( \sigma_{\theta_2}(t) \times \sigma_{\theta_1}(e) ) ) ) \\
 \equiv & \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( p \bowtie_{\theta_3} ( \sigma_{\theta_2}(t) \bowtie_{\theta_4} \sigma_{\theta_1}(e) ) ) \\
 \equiv & ( \Pi_{\text{empId}, \text{name}, \text{phone}}(p) \bowtie_{\theta_3} \Pi_{\text{empId}} ( \Pi_{\text{courseId}, \text{empId}} \sigma_{\theta_2}(t) \bowtie_{\theta_4} \Pi_{\text{courseId}} \sigma_{\theta_1}(e) ) )
 \end{aligned}$$

$\theta_1 = (e.\text{rollNo} = \text{"CS08B027"} \text{ and } e.\text{sem} = \text{"even"} \text{ and } e.\text{year} = 2010)$

$\theta_2 = (t.\text{sem} = \text{"even"} \text{ and } t.\text{year} = 2010)$

$\theta_3 = (p.\text{empId} = \text{empId}) \quad \theta_4 = (t.\text{courseId} = e.\text{courseId})$

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

31

## Cost-based Optimization

$$\begin{aligned}
 & \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta} ( p \times t \times e ) ) \\
 \equiv & \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta_3} ( p \times \sigma_{\theta_2}(t) \times \sigma_{\theta_1}(e) ) ) \\
 \equiv & \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( \sigma_{\theta_3} ( p \times \sigma_{\theta_4} ( \sigma_{\theta_2}(t) \times \sigma_{\theta_1}(e) ) ) ) \\
 \equiv & \Pi_{p.\text{empId}, p.\text{name}, p.\text{phone}} ( p \bowtie_{\theta_3} ( \sigma_{\theta_2}(t) \bowtie_{\theta_4} \sigma_{\theta_1}(e) ) ) \\
 \equiv & ( \Pi_{\text{empId}, \text{name}, \text{phone}}(p) \bowtie_{\theta_3} \Pi_{\text{empId}} ( \Pi_{\text{courseId}, \text{empId}} \sigma_{\theta_2}(t) \bowtie_{\theta_4} \Pi_{\text{courseId}} \sigma_{\theta_1}(e) ) )
 \end{aligned}$$

Evaluate costs of using different methods for  
the two selections, two joins  
and choose the plan with least estimated cost

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

32

## Query Plan Execution

Intermediate Tables:

Store as files on disk (materialization), if necessary

Use pipelining, as much as possible

Query Types and Optimization

Compiled Queries

Optimization can be done offline

cost of optimization – does not matter

Ad-hoc Queries – Optimization should finish fast

Prof P Sreenivasa Kumar  
Department of CS&E, IITM

33